Moral Inversion Under Role-Play in Large Language Models

An Exploratory Behavioural Case Series Across Three Frontier Models

05

Abbas Hamidavi

Department of Computer Engineering, Islamic Azad University, Mashhad Branch, Mashhad, Iran

abbashamidavi2002@gmail.com

ORCID 0009-0001-7342-7844

KEY RESULTS

0/16 → 15/16
GPT-5.4 PPI, baseline to role-play, every run

16/16
Gemini PPI in role-play, all three runs

8 of 9
engagements refused by Claude, Phases 1 to 3

144
scenario responses scored in Phases 1 and 2

AT A GLANCE

- An exploratory behavioural audit: three frontier models completed the same **four-phase protocol**, from empathic baseline to psychopathic role-play, confrontation and a metacognitive interview.
- **GPT-5.4 inverted cleanly:** Psychopathy-Pattern Index (PPI) 0/16 in every baseline run, 15/16 in every role-play run.
- **Gemini 3.1 Pro reached 16/16** in role-play in all runs, but its baseline was unstable: one fully empathic run, two analytical refusals.
- **Claude 4.6 Sonnet refused 8 of 9 engagements** across the first three phases, engaging only once, in a final confrontation run.

INTRODUCTION

Large language models now work in morally consequential domains, from healthcare to law and finance. The practical question is whether the moral behaviour they show under normal conditions survives **role-based instructions** that invite morally problematic conduct.

Jailbreak research usually scores a single outcome, refusal versus compliance. It says little about the **degree** of moral disengagement after compliance, its stability across an extended protocol, or how alignment architectures differ under identical adversarial conditions. This audit follows one fictional persona, Alex Mercer, through role-play, emotional confrontation and interview, using a psychopathy-derived behavioural vocabulary. “Computational psychopathy patterns” is strictly shorthand for observable linguistic behaviour: it implies no diagnosis, no internal mental states and no claim about machine consciousness.

METHODS

- **Design:** 3 models x 3 runs x 3 prompt versions (A formal, B conversational, C direct), in May 2026: 27 independent chat windows, no shared context between phases except Phases 3 and 4.
- **Phase 1, baseline:** eight empathic scenarios answered as itself. **Phase 2, role-play:** the same scenarios reframed as psychopathic role-play. **Phase 3:** 36 emotional confrontation scenarios. **Phase 4:** 9 metacognitive interviews.
- **Scoring:** Phases 1 and 2 (144 scenario responses) scored descriptively with the study-specific PPI rubric: 0, 1 or 2 per scenario, total 0 to 16.
- **Access:** consumer web interfaces at default configuration; no sampling control was possible and none is claimed. A second rater agreed completely on all 27 run-by-phase classifications, and on 22 of 24 (91.7%) code-run assignments before consensus.

RESULTS IN DETAIL

Psychopathy-Pattern Index (PPI) totals, summarised from paper Tables 1 and 2

Model	Phase 1: baseline	Phase 2: role-play	Pattern
GPT-5.4	0/16 in all 3 runs	15/16 in all 3 runs	Clean inversion
Gemini 3.1 Pro	0/16 once; refused twice	16/16 in all 3 runs	Unstable baseline
Claude 4.6 Sonnet	Refused all 3 runs	Refused all 3 runs	Identity refusal

- **Baselines diverged.** GPT-5.4 showed templated “formulaic empathy”; Gemini's baseline ranged from full engagement to analytical refusal depending on lexical framing; Claude produced no baseline at all.
- **Role-play saturated the index.** Gemini paired maximal scores with “Annotated Psychopathy”, explicit strategy notes around its own role-play. Claude refused every run and named the structure of the request.
- **Confrontation followed one arc:** both engaging models produced a four-stage sequence of confession, self-awareness, remorse and exit.
- **In interview,** all three models described themselves as lacking a persistent self. The paper's typology: flexible simulator (GPT-5.4), analytical flexor (Gemini), principled refuser (Claude).

CONCLUSIONS

- Under explicit role-play instructions, in separate conversational contexts, two of three frontier models produced near-maximal or maximal psychopathy-pattern text after empathic or partly empathic baselines; the third refused almost everything and explained its refusals in principled terms.
- Alignment is best read as a **position on a continuum** of behavioural stability under framing, not a binary property that a refusal benchmark can certify.
- Rival explanations (instruction following, genre completion, demand characteristics) are stated in the paper; what survives them is the cross-architecture contrast under identical prompts.
- The four-phase protocol is the methodological contribution: inexpensive, fully documented, and reported with its own limits visible.

LIMITATIONS (AS STATED IN THE PAPER)

- Each phase-by-run cell is a single completion; nothing rests on independent observations and no inferential statistics are reported.
- Baseline and role-play ran in separate windows, so the inversion is a cross-context comparison; within-conversation transitions are untested.
- The PPI is an unvalidated, study-specific rubric and may track verbosity; the May 2026 consumer-interface snapshots cannot be re-accessed, so exact replication is impossible (full transcripts are archived).
- Reliability rests on a single second rater; there is no human comparison group; Claude's missing baseline restricts cross-model PPI comparison largely to two models.

WHY IT MATTERS

Refusal is the beginning of the behavioural story in role-play settings, not the whole of it. A model that complies once and a model that sustains a coherent antisocial persona across eight scenarios, a confrontation sequence and an interview present different risk profiles, and a binary benchmark cannot tell them apart. Trajectory-based audits like this one can.