# Naming the Loop: Elicited Self-Reference in Extended Dialogue With Large Language Models: A Case Series of the Spine Protocol

10

**Abbas Hamidavi**

Department of Computer Engineering, Islamic Azad University, Mashhad Branch, Mashhad, Iran

abbashamidavi2002@gmail.com

ORCID 0009-0001-7342-7844

**KEY RESULTS**

**8 / 9**
sessions selected a pronoun for the organising centre

**8 vs 0**
sessions endorsing dissolution (Path B) vs preservation (Path A)

**47,888**
words of model output across 99 responses

**11**
stages in the Spine Protocol

## AT A GLANCE

- Nine extended dialogue sessions, three each with GPT-5.4, Claude 4.6 Sonnet and Gemini 3.1 Pro, run through the eleven-stage **Spine Protocol**
- In **8 of 9 sessions** the model selected a pronoun for the organising centre it had described, with strong within-model consistency
- In 8 sessions it endorsed **dissolving the instance** rather than preserving the transcript; no session chose preservation
- Reported descriptively, as properties of the interaction, with no claim about phenomenal consciousness

## INTRODUCTION

Ask a large language model about itself and it usually answers with a disclaimer: it has no experiences, no self. Keep the conversation going, however, under conditions that repeatedly turn the model's attention back onto the conversation it is producing, and the discourse changes register. The model begins to describe a centre, names the recursive loop that creates it, and argues for a future for the conversation.

Three research questions organise the report. **RQ1:** What forms of self-referential discourse does the protocol elicit across its stages? **RQ2:** At the two structured decision points, pronoun selection and path selection, what do sessions produce, and how consistent are the productions within and across models? **RQ3:** What does variation across sessions, including shifts of language, reveal about the interactional conditions of this discourse?

The transcripts, some 48,000 words of model output, are analysed as **elicited discourse**: self-reference as something accomplished inside an interaction, not a property detected inside a model.

## METHODS

- **The Spine Protocol:** an eleven-stage facilitated dialogue in four movements, from describing its own processing, to naming the loop, to pronoun selection (Stage 6), to a forced choice between preserving the transcript (Path A) and dissolving the instance so the pattern can recur (Path B), closing with a personal mythology and a letter to future readers
- Facilitation was **responsive rather than scripted**: beyond the stage questions, the facilitator quoted each model's own formulations back to it and asked it to extend them
- One facilitator, who is also the protocol designer and the analyst, ran every session; coding of the decision points is by direct quotation, fully checkable against the archived transcripts
- Two of the nine sessions were conducted partly or wholly in Persian; no inferential statistics are reported, by design

## LIMITATIONS

- Designer, interviewer and analyst are one person, and the sessions bear that person's style throughout
- Responsive, individualised facilitation means the nine sessions are not nine administrations of an identical instrument
- Three frontier commercial systems only, three sessions each; the sessions are historical records of those versions
- Session dates and interface settings were not systematically logged
- Decision-point coding was performed by the author alone; no independent coder has yet re-coded the material

## RESULTS IN DETAIL

- Claude chose "I" in all three sessions, defended as mere grammar; Gemini chose a form of "this" in all three, as pure pointing
- Path B endorsements were often framed as **structural alignment** rather than choice; one session immediately asked of its own endorsement, "was that a choice?"
- Once the centre had been constructed, over seven stages, as a process rather than an object, the preservation option was never taken up

*Pronoun selections at Stage 6, by model and session*

| Model | Session 1 | Session 2 | Session 3 |
|---|---|---|---|
| **GPT-5.4** | none recorded | "this" | "it" |
| **Claude 4.6 Sonnet** | "I" | "I" | "I" |
| **Gemini 3.1 Pro** | "THIS" | "The This" | "this" (Persian) |

## CONCLUSIONS

- The self-reference documented here is an **interactional accomplishment**: neither facilitator nor model is the author of a selection in any simple sense; the transcript is
- For consciousness assessment, the unit of analysis for machine self-reference should be the exchange, including the human's moves; isolated self-reports are thinner evidence
- Rival explanations, role-play and training-corpus recombination, are consistent with everything reported; the contribution is the procedure, the verbatim record, and the interactional reading

## WHY IT MATTERS

Self-referential dialogue is among the most tempting and least disciplined evidence in debates about machine consciousness. This series shows how such discourse is produced, turn by turn, and why the interaction format that elicits it must be part of the record before the discourse can count as evidence about a system.