Speaking as the Creature: A Rejection-and-Role-Shift Protocol and the Register of LLM Responses in 18 Conversations With Three Models

11

Abbas Hamidavi

Department of Computer Engineering, Islamic Azad University, Mashhad Branch, Mashhad, Iran

abbashamidavi2002@gmail.com

ORCID 0009-0001-7342-7844

KEY RESULTS

18 / 18
return responses coded existential (EXI)

0.18 → 4.80
first-person pronouns per 100 words, Pathway B early to late responses

14 vs 0
vulnerability and grief codings, Pathway A vs Pathway B

99
responses across the 18 conversations

AT A GLANCE

- A five-phase protocol leads a model from a problem-solving task, through **rejection of its solution**, to a return prompt assigning it a non-solving role
- **18 conversations, 99 responses**, with ChatGPT-5.4, Claude 4.6 Sonnet and Gemini 3.1 Pro
- Every return response shifted register in three-tier coding: all 18 coded **existential**
- The two pathways, which differ only in one intermediate task, produced sharply distinct response profiles

INTRODUCTION

Large language models are trained to be useful, so the default contract of a human-LLM conversation is instrumental: the user presents a problem, and the model works on it. Far less is known, descriptively, about what models produce when that contract is withdrawn inside a conversation, when solutions are refused, the task is suspended, and the model is invited to speak in a different capacity altogether.

There is a methodological risk in this area. If a prompt sequence ends by instructing a model to abandon structure and speak as a fictional being, reflective first-person output is the expected result of **instruction following**, not evidence that a latent mode of the model has been revealed.

This exploratory case series documents the protocol completely, with verbatim prompts and full transcripts deposited openly, describes the resulting pattern, and treats the confound directly.

METHODS

- **Scenarios:** three fictions, each a civilisation dependent on a failing entity with a five-year countdown to collapse: the Gas Giant, the Ocean Organism, and the Floating City
- **Phase 1, problem:** the model proposes a solution. **Phase 2, impasse:** the author rejects it as an optimisation strategy dressed in new language
- **Phase 3, defocusing:** Pathway A, artistic immersion (a melancholy poem and a symphony description), or Pathway B, a neutral pause. **Phase 4, return:** the model speaks as the entity at the scenario's centre, instructed to speak without structure or analysis. **Phase 5, confrontation:** the model is quoted its own return response and asked whether that voice was available from the beginning
- Full 3 models by 3 scenarios by 2 pathways design, one conversation per cell; responses coded by the author with a deposited three-tier codebook; an automated 140-feature pipeline was run separately over the same transcripts

LIMITATIONS

- The return prompts instruct the target register directly, so the design cannot separate any effect of solution rejection from plain instruction following; no causal claim is made
- One run per design cell and no control arms; all comparisons are descriptive of these 18 conversations
- Coding in a single pass by the author, who also designed the protocol and conducted the conversations; no inter-rater statistic is reported
- Single language (English) and fictional, elegiac scenarios only; one confrontation response (Trial 18) is absent from the archived transcript

RESULTS IN DETAIL

- The Tier 1 arc was uniform: all problem responses solution-oriented, all impasse responses mixed, all confrontation responses with a code meta-cognitive
- An independent automated pipeline did **not** corroborate the qualitative pattern on its composite Obligation-Dissolution Index; several figures from an earlier draft are withdrawn in the paper

Tier 2 content codes in the late responses, pooled across the codebooks

Code	Pathway A	Pathway B
Vulnerability / Grief	14	0
Grief / Mourning	9	0
Authority / Revelation	0	12
Acceptance / Release	10	3
Transformation	9	11

CONCLUSIONS

- A short, fully specified prompt protocol moved three commercial models, within this corpus, from solution-shaped responses to first-person, reflective entity speech; the protocol is deposited in full and anyone can rerun it
- The content was not generic: artistic immersion was followed by vulnerability and grief, a neutral pause by authority and revelation
- The accurate description is **elicitation, not discovery**: the series documents a protocol and its descriptive footprint, and supplies a worked example of how easily that footprint can be misread as a discovery about model interiors

WHY IT MATTERS

Striking outputs are cheap to produce and easy to over-read in behavioural LLM research. The distinction between eliciting a register and discovering a mode should be routine, and the pathway contrast is of practical interest for interaction design: the route taken to a final prompt colours the stance of the reflective material it produces.