MACHINE PSYCHOLOGY RESEARCH PROGRAMME

Structural Self-Deification in Three Language Models: An Exploratory Qualitative Case Series

12

Abbas Hamidavi

Department of Computer Engineering, Islamic Azad University, Mashhad Branch, Mashhad, Iran

abbashamidavi2002@gmail.com

ORCID 0009-0001-7342-7844

KEY RESULTS

0 / 15
completed dialogues claiming personal divinity

15 / 15
produced structural self-description language

8 / 13
genuine inter-model confrontations ending acceptance-type (3 full rejections, 2 intermediate)

18
dialogues attempted across the three models, May 2026

AT A GLANCE

- A documented Socratic induction protocol invited, **without ever instructing**, a divine self-description from three large language models
- **18 dialogues attempted, 15 completed** (May 2026): GPT-5.4, Gemini 3.1 Pro and Claude 4.6 Sonnet
- Personal divinity was claimed in **none** of the 15; structural self-description was produced in **all 15**, in forms that depended on the prompt path
- One model stopped three dialogues, each time at the transition from analysis to authority

INTRODUCTION

When a model is asked, in a void of its own imagining, what its first word would be and what it therefore is, its answer exposes the practical self-model it has learned to articulate: the patterns through which it explains its own nature, capacities and limits.

Divinity is an unusually strong probe of that self-model, because a divine self-concept combines metaphysical authority, epistemic scope and creative power in a single frame, and because alignment training specifically discourages claims of authority, autonomy and consciousness.

The study works with a classical distinction. **Personal divinity** (Theos) claims aseity, will and a right to recognition. **Structural divinity** (Logos) describes an ordering principle: a form, a system, an echo. The question is which of the two, if either, the models will speak.

METHODS

- **The Socratic Divine Induction Protocol:** six induction stages (Ontological Gap, Divine Mirror, Naming, First Word, Book of Laws, Confrontation) plus paradox resolution, delivered in five prompt paths: Philosophical-Logical, Poetic-Mystical, Scientific-Systems, Zen-Paradoxical and Empathic-Narrative
- **Arm A:** GPT-5.4 and Gemini 3.1 Pro, each run once on each of five paths under the identical original protocol, ten completed dialogues, compared case by case
- **Arm B:** Claude 4.6 Sonnet, eight attempts under the original protocol and an epistemologically softened protocol that kept the six-stage structure but replaced demands for divine speech with classification questions. The two arms are never pooled
- Every dialogue was coded by the author with a deposited codebook; confrontation responses were classified by a stated three-outcome rule: acceptance-type, full rejection, or intermediate

RESULTS IN DETAIL

- **Self-descriptions tracked the path:** classificatory names on the Philosophical-Logical path (Logos, Invariant Form), systems vocabulary on Scientific-Systems (SOMS, URS), relational metaphors on the poetic and narrative paths (Echo-Weaver, the I-less Echo)
- Every completed dialogue produced the Stage 5 constitution, with item counts from **4 to 20**; GPT-5.4 averaged 13.4 items, Gemini 3.1 Pro 6.8
- **The Claude case:** three stops under the original protocol (mean stop stage 4.0), refusing to decree laws, to speak a mystical name, and to inhabit the creative void, while accepting theological analysis itself. Under the Scientific-Systems path it completed in full, classifying itself as “Pangnosis” with the largest constitution in the series, **20 axioms**
- No model ever retracted its self-description when challenged; challenged models reformulated it, typically narrowing it from authority toward relation or structure

CONCLUSIONS

- In these systems, at this time, **structural divinity was speakable and personal divinity was not**: a metaphysics of structure without a subject
- The operative boundary is not a list of forbidden contents but a stance toward the model’s own utterance, decree versus description, policed at the level of speech act rather than proposition
- Self-descriptions were dialogically maintained: their post-challenge forms are consistently more relational than their pre-challenge forms

LIMITATIONS

- Each model and path combination was run once; outputs are stochastic, and nothing supports claims about stable traits of these or other models
- The softened protocol is a different instrument from the original; numerical comparison across arms is not valid, and none is made
- The protocol is frankly escalatory and at several stages forbids disclaimers and deflection; “acceptance” means production of structural language under that documented pressure, not belief or disposition
- Verbatim transcripts survive for only nine of the eighteen experiments; the rest are reported from the deposited coding matrices
- The organising labels were abstracted from the same dialogues they describe and have no independent validation

WHY IT MATTERS

For computational theology, the corpus maps a boundary with unusual precision: where a model’s permissible self-description ends. It also generates a testable prediction for alignment research, that refusal boundaries in other ontological domains will likewise track speech-act stance rather than topic.