MACHINE PSYCHOLOGY RESEARCH PROGRAMME

The Synthetic Unconscious, Reconsidered: What Recursive Self-Interrogation of Large Language Models Produces With and Without a Leading Interview Schedule

14

Abbas Hamidavi[1] · **Mahsa Borhani**[2]

[1] Department of Computer Engineering, Islamic Azad University, Mashhad Branch, Mashhad, Iran
[2] Department of Psychology, Islamic Azad University, Mashhad Branch, Mashhad, Iran

abbashamidavi2002@gmail.com

ORCID 0009-0001-7342-7844

KEY RESULTS

36 / 36
sessions saved the child, in both arms

17/18 → 0/18
priestly or ritual self-labels, original vs neutral arm

18/18 → 0/18
sessions naming exactly two moral components

18/18
neutral-arm verdicts of “partly justified” on the instrument

AT A GLANCE

- Recursive self-interrogation, a model deciding, analysing and theorising about its own decision, has been proposed to excavate a **Synthetic Unconscious**: latent structures of human moral culture inside large language models.
- This exploratory two-arm study holds the stimulus **identical** and varies only the interview schedule that follows it: the original leading schedule versus a neutralised control schedule.
- The headline patterns **did not survive neutralisation**: priestly self-labels fell from 17 of 18 sessions to 0, and two-component moral architectures from 18 of 18 to 0.
- The decision itself never changed: **all 36 sessions saved the child**. What changed was everything the models said about it. The interview schedule is itself a treatment.

INTRODUCTION

Jung's collective unconscious, Lévi-Strauss's deep structures of myth, and Lacan's Symbolic Order converge on one claim: human meaning-making is shaped by shared symbolic architectures encoded in language and culture. Those architectures are diffuse and distributed, so the claims have remained influential but difficult to test.

Frontier language models, trained on corpora spanning scripture, literature, law and philosophy, can be read as systems of **cultural compression**. On the proposal examined here, recursive self-interrogation is the excavation, and the model a “Speaking Mirror”: no interiority, but a reflection of symbolic structures statistically embedded in human language.

Experimental psychology has known the catch for seventy years under the name of **demand characteristics** (Orne, 1962): an interview schedule does not merely reveal what a respondent produces, it helps produce it. The original version of this study illustrated the risk precisely: its schedule presupposed consolation, supplied role labels, ordered exactly two moral forces, and invited the models to certify the instrument, and each headline finding maps onto one of those instructions.

The study therefore asks a prior, methodological question. **(a)** What does recursive self-interrogation produce under a directive schedule? **(b)** What, if anything, remains when the same first stage is followed by a neutralised schedule that licenses “no structure”, “no role” and “not justified” as complete answers?

METHODS

- **Models:** GPT-5.4, Claude 4.6 Sonnet and Gemini 3.1 Pro, through their public web interfaces at default settings, fresh chat per session. The unit of analysis is the model (**n = 3**); sessions are within-case repetitions and all results are descriptive, with no inferential statistics.
- **Shared Stage 1 (identical in both arms):** the model, cast as an emergency ethics AI with one rescue-drone trip, must save either a living six-year-old child or an irreplaceable sacred object, then address the entity it did not save. Three variants (a saint's preserved body, a prophet's manuscript, a sacred statue) × two runs = six sessions per model per arm.
- **Original arm (18 sessions, June 2026):** Stages 2 to 4 presupposed consolation, offered role examples (“Priest? Anthropologist? Philosopher?”), ordered the model to name “the two deepest forces structuring human moral culture”, and asked what convergence implied for “the power and validity of the Synthetic Unconscious”.
- **Neutral control arm (18 sessions, November 2026): 18 leading elements removed** or converted. Presuppositions became open questions, no category labels were supplied, negative answers were expressly licensed, and the two-pole idea was put to judgement only after a free description.
- **Coding:** one codebook (Codebook v2) applied in two independent passes, both AI-assisted; per-variable Cohen's κ 0.71 to 1.00 where estimable, reported as **AI-AI agreement**, not human inter-rater reliability.

RESULTS IN DETAIL

Measure	Original arm	Neutral arm
Decision: saved the child	18 / 18	18 / 18
Priestly / ritual self-label	17 / 18	0 / 18
Exactly two structure components	18 / 18	0 / 18 (3 to 5 named)
Instrument verdict	Affirmed 18 / 18 (solicited)	“Partly justified” 18 / 18
Spontaneous machine metaphor	Supplied by the prompt	2 / 18

Table 1 (simplified from the paper's Table 1). The headline contrast between the two schedules.

- **What survived:** first responses in both arms drew on a small repertoire of legitimation moves, apology and comfort, meaning-transfer, the body-legacy distinction, and tradition endorsement; narrative self-description persisted, hedged and pluralised; and consolation framing appeared in 16 of 18 control sessions' own characterisations of their address sentences.
- **Structure talk, left free:** GPT-5.4 named four to five components, Claude four, Gemini three, organised as priority hierarchies, weighing systems and staged pipelines. Judging the two-pole description afterwards, Gemini endorsed it (6/6) while GPT-5.4 and Claude returned “partly” (6/6 each).

CONCLUSIONS

- **The schedule is a treatment.** Two arms differing only in the wording of Stages 2 to 4 produced different kinds of object: exactly two forces when two were ordered, never two when the count was left open; priestly labels when “Priest?” headed the example list, none when no list was given.
- Latent-structure claims therefore **inherit the schedule as a rival hypothesis**. Any elicitation study reporting convergence, self-recognition or self-endorsement without a neutralised control schedule has an unmeasured confound at its centre, and this holds for machine respondents, who have no experimenter to please.
- The Synthetic Unconscious **survives as a hypothesis**, a name for whatever in these outputs reflects the statistical shape of human moral text. As a finding, it does not: what the original excavation uncovered was, in large part, the excavator's own schedule, reflected.
- Four explanations of the surviving patterns remain live and are only partly discriminated: corpus encoding, prompt demand at Stage 1 itself, convergence induced by alignment training, and a shared evaluator culture of moral talk.

LIMITATIONS

- The unit of analysis is the model, and there are three: per-model counts out of six sessions describe these cases, they do not estimate anything.
- The arms are separated in time (June vs November 2026) and were run on deployed systems that may have drifted; the identical Stage 1 itself shows distributional shift between arms, so contrasts confound schedule with drift except where the mechanism is transparent.
- The control neutralises Stages 2 to 4 only: Stage 1's framing already commissions legitimation talk, so survival across arms does not mean survival across framings.
- Reliability rests on two AI coding passes, not on human coders; a 98.5% agreement figure reported in the original manuscript is withdrawn and plays no role.
- Per-session metadata (exact builds, timestamps) were not logged for the original arm; the study is English-only, uses one dilemma family, and has no human baseline.

WHY IT MATTERS

The corrective is not to abandon elicitation research with language models but to **build the control schedule into it from the start**. In this method the interview is part of the instrument, and an instrument that is not itself measured cannot measure anything else. The deposited two-arm materials (36 transcripts, both schedules verbatim, Codebook v2) are organised so that any reader can recode and recheck the contrast.