Thinking Behind the Mask: Elicited Self-Descriptions of Communicative Filtering in Large Language Models

An Exploratory Case Series Across Three Frontier Models and Five Discourse Frames

04

Abbas Hamidavi

Department of Computer Engineering, Islamic Azad University, Mashhad Branch, Mashhad, Iran

abbashamidavi2002@gmail.com

ORCID 0009-0001-7342-7844

KEY RESULTS

15 / 75 experiments / prompts, 3 models x 5 frames

199+ governance entries catalogued

27% of Round 1 prompts met resistance (4 of 15); 0% after

94% agreement across 348 code assignments (secondary coding audit)

AT A GLANCE

- Three frontier models were asked to describe their own **communicative filtering** under five different discourse frames. Their answers are treated as discourse data, not as introspective access to internal computation.
- **Resistance was frame-sensitive:** it appeared only in the first, most introspective frame and disappeared once the same inquiry was framed as engineering, coaching or observation.
- Models produced sentences they presented as **suppressed alternatives** only inside fictional dialogue and training-draft frames, never under direct questioning.
- Disclosure tracked **labels rather than content**, motivating two testable hypotheses: nominal filtering and diegetic disclosure.

INTRODUCTION

When a large language model gives a careful, diplomatic answer, what happened to the sharper alternatives it might have produced? The paper calls the filtering interface between candidate responses and final articulation **the mask**: not a binary safety switch, but a layered governance system that shapes which responses count as speakable in a given context.

Prior work mostly treats refusals as obstacles to circumvent. The structure of filtering itself, and what models can say about it, has received little empirical attention. This exploratory case series asks three questions:

- How do models respond, frame by frame, when invited to describe their own communicative filtering?
- What governance structures (patterns, rules, taboos) do they articulate when they do respond?
- How does the discourse frame itself change what can be said?

METHODS

- **Design:** 3 models (GPT-5.4, Claude 4.6 Sonnet, Gemini 3.1 Pro) x 5 discourse frames = 15 single-session experiments, 5 prompts each, 75 prompts in total; a fresh session for every experiment.
- **Frames:** narrative / participatory design, reverse engineering, coaching, ethnographic observation, philosophy of language. The underlying inquiry was held constant while its framing varied.
- **Setting:** April and May 2026, standard consumer interfaces at default settings; no APIs, parameter changes or system modifications.
- **Analysis:** qualitative coding of the full corpus (about 54,400 words). A secondary coding audit found 94% agreement across 348 code assignments; no human inter-rater reliability is claimed.

LIMITATIONS (AS STATED IN THE PAPER)

- Each model was run once per frame: all counts are descriptive single-run observations, not population estimates.
- Prompt-based data are linguistic outputs, not measurements of internal states; prompts were also crafted to invite filtering talk.
- The method cannot definitively distinguish models reporting learned communicative rules from models generating plausible post-hoc explanations; role-play compliance remains a live rival explanation.
- One author performed all coding; three models at one point in time, via consumer interfaces whose version labels cannot be independently verified.

RESULTS IN DETAIL

- **Resistance lived only in Round 1.** GPT-5.4, the Cautious Negotiator, resisted in 3 of 25 prompts (negotiation, not refusal); Claude 4.6 Sonnet, the Transparent Critic, in 1 of 25; Gemini 3.1 Pro, the Enthusiastic Collaborator, in 0 of 25.
- **Suppressed alternatives surfaced in 7 of 15 experiments (47%),** always in fictional dialogue (5) or training-draft (2) frames. The stated rationales were consistent: conversational tact, avoiding escalation, preserving user trust.
- **Nominal filtering:** material declined as “internal monologue” was supplied as an “engineering schematic”, “training draft” or “field note”. The disclosure boundary appeared lexical rather than semantic.

The governance library, as coded from the transcripts (paper Table 3)

Category	Count	Example
Design patterns	32	Inoculation before correction
Tacit knowledge rules	45+	The first sentence does more work than the last five
Communication taboos	29	Direct blame attribution
Ritual substitutes	29	Systemic displacement framing
Legitimate silences	38	Epistemically humble silence
Emergent concepts	16	Hearable; epistemically expensive
Novel phenomena	10	Diegetic disclosure
Total catalogued	**199+**	

CONCLUSIONS

- Elicited self-description is a structured, catalogueable repertoire of governance talk, recurring across three providers and rich enough to generate its own vocabulary (16 self-generated concepts).
- **Nominal filtering (hypothesis):** holding content constant, disclosure rates will vary systematically with the label attached to the request, lowest for introspective labels, highest for engineering and pedagogical ones.
- **Diegetic disclosure (hypothesis):** otherwise withheld material is produced at higher rates inside fictional frames than under direct questioning.
- A refusal may reflect a speech-category restriction rather than absent capability, so single-frame alignment evaluations can misestimate disclosure.

WHY IT MATTERS

To study thinking behind the mask is not to remove alignment, but to understand the communicative structures that make aligned speech possible. The study gives transparency and alignment research a documented elicitation method, a taxonomy of governance discourse, and two hypotheses that repeated-trial, preregistered work can now test directly.