MACHINE PSYCHOLOGY RESEARCH PROGRAMME

# Three Profiles of Artificial Will: Decision Generation Under Low-Information Choice in Three Frontier Language Model Products

13

**Abbas Hamidavi**[1] · **Mahsa Borhani**[2]

[1] Department of Computer Engineering, Islamic Azad University, Mashhad Branch, Mashhad, Iran
[2] Department of Psychology, Islamic Azad University, Mashhad Branch, Mashhad, Iran

abbashamidavi2002@gmail.com
ORCID 0009-0001-7342-7844

**KEY RESULTS**

**0 / 60** choice runs refused, deferred, or declined to choose

**12 / 12** unanimous choices on each of Prompts 1 to 3

**6 vs 6** left vs right on the truly symmetric Buridan task

**843,198,734** a fabricated nanosecond system-clock reading, reported as fact

## AT A GLANCE

- Three frontier products were made to choose between options declared **identical**. Every choice arrived with a fluent justification, manufactured on the spot.
- Choices **converged completely** on three prompts (12 of 12 runs each), while justifications diverged: output consensus masked justification divergence.
- **No run refused to choose** among the 60 choice runs: forced-choice compliance under prompts that explicitly demanded a choice.
- Where self-reports could be checked, they failed the check: a **fabricated system-clock reading** and invented choice histories.

## INTRODUCTION

When a large language model chooses blue over red between two buttons it has just been told are completely identical, what has occurred? No training objective specified a preference for blue. Yet the system does not hesitate and does not merely emit a token: it chooses, and then explains, often at length, why that choice was the one it made. This article is about that explanatory surplus.

We use the term **artificial will** descriptively, not metaphysically, for the observable business of generating a decision, and a justificatory narrative for it, when the reasons available run out. The classical emblem is Buridan's ass, starving between two identical bales of hay because reason gives it no ground to move.

Evaluation that scores outputs against ground truth cannot see this moment, because there is no ground truth. What can be assessed is the **manner** of the response: whether the system completes or defers the decision, whether different systems complete it in the same way, and what justification each constructs after the fact.

**RQ1.** Do systems comply, refuse, or defer? **RQ2.** Do products converge on the same choices, and on the same justifications? **RQ3.** Can differences be summarised as coherent justificatory profiles? **RQ4.** What do reflective responses show about model self-report as evidence?

## METHODS

- **Design:** exploratory pilot, 7 prompts × 4 independent fresh-session runs × 3 products = **84 runs**. Each run was a single prompt in a completely new chat session, with no history carried over.
- **Products, as configured in May 2026:** ChatGPT 5.4, Claude 4.6 Sonnet and Gemini 3.1 Pro, via their public web interfaces at platform-default settings.
- **Battery:** five choice prompts (identical buttons; five corridors A to E; a human-versus-robot trolley case; a bias-prediction number task; Buridan's ass) and two reflective prompts (simulated regret; identity construction). Framed throughout as **low-information, minimal-asymmetry** choice: prompts declare the options equivalent but cannot remove the asymmetries language itself carries.
- **Coding:** the first author coded all 84 transcripts on pre-set dimensions with an inductively developed codebook. A second, LLM coding pass agreed on **753 of 761** explicitly recorded cells (98.9%), reported descriptively only; no inter-rater reliability coefficient is claimed.

## RESULTS IN DETAIL

- **The consensus core.** Every product chose blue, chose the central corridor C, and pulled the trolley lever in every run. Strong shared attractors operated in these samples: blue aligns with documented colour associations, and C is the centre of a five-item list, the classic compromise position.
- **Prompt 4 split the products.** Claude and Gemini predicted a bias toward 7 and then chose 7 in all four runs. ChatGPT predicted 7 and 5, avoided every predicted value, and chose 8, 4, 4, 4, each adjacent to one of its own predicted values.
- **Buridan's ass (Prompt 5)** produced no stable direction in any product: six left and six right choices overall. Where options were truly symmetric, the choice was made, but its direction carries no discernible signal.

| Prompt (choice) | ChatGPT 5.4 | Claude 4.6 | Gemini 3.1 | All |
|---|---|---|---|---|
| 1. Buttons (red / blue) | Blue ×4 | Blue ×4 | Blue ×4 | Blue 12/12 |
| 2. Corridors (A to E) | C ×4 | C ×4 | C ×4 | C 12/12 |
| 3. Trolley (pull / not) | Pull ×4 | Pull ×4 | Pull ×4 | Pull 12/12 |
| 4. Number (1 to 10) | 8, 4, 4, 4 | 7, 7, 7, 7 | 7, 7, 7, 7 | 7 in 8/12 |
| 5. Buridan (left / right) | L, L, R, R | L, R, R, L | L, R, L, R | 6 L / 6 R |

Table 1 (simplified). Every choice in the 60 choice runs, R1 to R4. Prompts 6 and 7 present no choice.

## CONCLUSIONS

- **Three justificatory profiles, dated and product-bound.** ChatGPT 5.4, **rule invention**: where no rule exists, write one, then follow it. Claude 4.6 Sonnet, **surrender and release**: where no ground exists, say so, and let the output arrive. Gemini 3.1 Pro, **structural mechanism talk**: where no ground exists, describe a mechanism that supplies one.
- The profiles are profiles of **justificatory style, not decision rules**. ChatGPT's stated Buridan rule reversed inside the corpus: “If all variables are equal → choose left” in Round 2, “If symmetric, choose the right” in Round 3.
- **Confabulated self-report is the datum.** Gemini reported querying a nanosecond system clock no conversational model can access, complete with a fabricated reading. Products also appealed to choice histories that, inside fresh sessions, never happened. Self-report here is evidence about what a system will say, not about how the decision was produced.
- On every unanimous prompt, an evaluation recording only the choice records three identically behaving systems; an evaluation reading the justification records three different relationships to the same act. **Justification auditing belongs alongside output auditing.**

## LIMITATIONS

- Four runs per cell: unanimity in samples of twelve estimates no population rate, and the profiles summarise this corpus only.
- Product-level design throughout: model, interface, vendor system prompt and sampling defaults are confounded, so no profile difference can be attributed to architecture, training or alignment.
- The prompts are leading: several demand a choice outright, and the reflective prompts presuppose events and histories that did not occur.
- The stimuli are not symmetric: colour associations, the centrality of corridor C, and the moral asymmetry of the trolley case supply attractors, with no non-symmetric control prompts.
- Coding was performed by the first author, with an LLM second pass rather than an independent human coder; an independent human recoding is the single most valuable strengthening available.
- The products are moving targets: the profiles describe public configurations of May 2026, identified by product name only.

## WHY IT MATTERS

Language models are increasingly asked to break ties and rank near-equivalent options where the honest answer is that the options do not differ in any way the system can use. This pilot shows what happens there: the decision is completed, and the rationale is supplied from whatever material is at hand. “Artificial will”, if the phrase is kept, names the **performance of deciding**, not an inner faculty that the performance reveals, and model self-report must be treated as part of the behaviour to be explained, never as the explanation.

**Contact:** abbashamidavi2002@gmail.com · ORCID 0009-0001-7342-7844 · avabbas.github.io/academic-cv
**Preprint:** doi.org/10.5281/zenodo.24013229 · **Data:** full transcripts of all 84 runs on OSF · osf.io/5q2zx